



東南大學
SOUTHEAST UNIVERSITY

OPERATING SYSTEM CONCEPTS

.....

Chapter 1. Introduction

A/Prof. Kai Dong



Contents

1. What Operating Systems Do
2. Computer-System Organization
3. Computer-System Architecture
4. Operating-System Structure
5. Operating-System Operations
6. Operating-System Functions
7. Kernel Data Structures
8. Open-Source Operating Systems



Warm-up

Discussion

- What Operating System(s) are you familiar with?
- What is the best Operating System in your opinion? And for what features?



Objectives



- To describe the basic organization of computer systems.
- To explain the evolution of operating system structures.
- To provide a grand tour of the major components of operating systems.
- To give an overview of the many types of computing environments.
- To explore several open-source operating systems.



Contents

1. What Operating Systems Do
2. Computer-System Organization
3. Computer-System Architecture
4. Operating-System Structure
5. Operating-System Operations
6. Operating-System Functions
7. Kernel Data Structures
8. Open-Source Operating Systems



What Operating Systems Do

What is an Operating System?

- A **program** that acts as an intermediary between a user of a computer and the computer hardware.
- Operating system goals:
 - Execute user programs and make solving user problems **easier**.
 - Make the computer system **convenient** to use.
 - Use the computer hardware in an **efficient** manner.



What Operating Systems Do

Computer System Structure

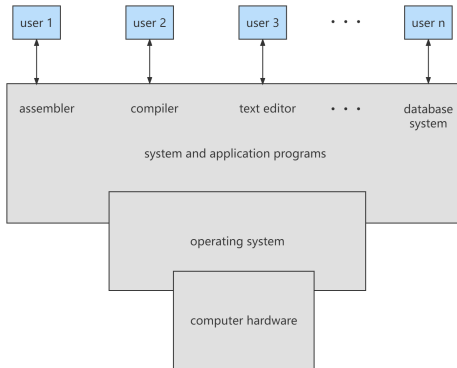
Computer system can be divided into four components:

- **Hardware** — provides basic computing resources.
 - CPU, memory, I/O devices
- **Operating system** — controls and coordinates use of hardware among various applications and users
- **Application programs** — define the ways in which the system resources are used to solve the computing problems of the users.
 - Word processors, compilers, web browsers, database systems, video games
- **Users**
 - People, machines, other computers



What Operating Systems Do

Computer System Structure (contd.)



- We can explore operating systems from two viewpoints: that of the **user** and that of the **system**.



What Operating Systems Do

User View

- Want **convenience, ease of use** and **good performance**.
- Don't care about **resource utilization**.
 - But *shared computer* such as *mainframe* or *minicomputer* must keep all users happy.
 - Users of *dedicate* systems such as *workstations* have dedicated resources but frequently use shared resources from *servers*.
 - Handheld computers are resource poor, optimized for usability and battery life.
 - Some computers have little or no user interface, such as embedded computers in devices and automobiles.



What Operating Systems Do

System View

- OS is a **resource allocator**.
 - Manages all resources.
 - Decides between conflicting requests for efficient and fair resource use.
- OS is a **control program**.
 - Controls execution of programs to prevent errors and improper use of the computer.



What Operating Systems Do

Defining Operating Systems

- No universally accepted definition.
- “Everything a vendor ships when you order an operating system”
 - is a good approximation,
 - but varies wildly.
- “The one program running at all times on the computer”
 - defines the kernel,
 - everything else is either
 - » a system program (ships with the operating system) , or
 - » an application program.



Contents

1. What Operating Systems Do
- 2. Computer-System Organization**
3. Computer-System Architecture
4. Operating-System Structure
5. Operating-System Operations
6. Operating-System Functions
7. Kernel Data Structures
8. Open-Source Operating Systems



Computer-System Organization

Computer Startup

- **Bootstrap program** is loaded at power-up or reboot.
 - Typically stored in ROM or EPROM, generally known as firmware.
 - Initializes all aspects of system.
 - Loads operating system kernel and starts execution.
- Think of:
 - What happens when a program runs / How instructions are executed?
 - Millions/billions of times a second,
 - » the processor fetches an instruction from memory,
 - » decodes it,
 - » executes it,
 - » moves on to the next till end.



Computer-System Organization

Von Neumann Model

- We have just described the basics of the **Von Neumann model** of computing.
- “...an electronic digital computer with parts consisting of
 - a processing unit containing an arithmetic logic unit and processor registers;
 - a control unit containing an instruction register and program counter;
 - a memory to store both data and instructions;
 - external mass storage; and
 - input and output mechanisms.”— Wikipedia “Von Neumann architecture”.




Computer-System Organization

Von Neumann Model (contd.)

The above Von Neumann Model defines/describes:

- How is a computer-system organized
 - One or more CPUs, device controllers connect through common bus providing access to shared memory.
- How is a computer-system operating
 - Concurrent execution of CPUs and devices competing for memory cycles.





Computer-System Organization

Computer-System Operation

- I/O devices and the CPU can execute **concurrently**.
- But how?
 - Each device controller is in charge of a particular device type.
 - Each device controller has a local buffer.
 - CPU moves data from/to main memory to/from local buffers.
 - I/O is from the device to local buffer of controller.
- By now communication is enabled, but how to cooperate?
 - Device controller informs CPU that it has finished its operation by causing an **interrupt**.



Computer-System Organization

Common Functions of Interrupts

- An interrupt is an input signal to the processor indicating an event that needs immediate attention.
- Interrupt transfers **control** to the interrupt service routine generally, through the **interrupt vector**, which contains the addresses of all the service routines.
- Interrupt architecture must save the address of the interrupted instruction.
- A **trap** or **exception** is a software-generated interrupt caused either by an error or a user request.
- An operating system is **interrupt driven**.

Computer-System Organization

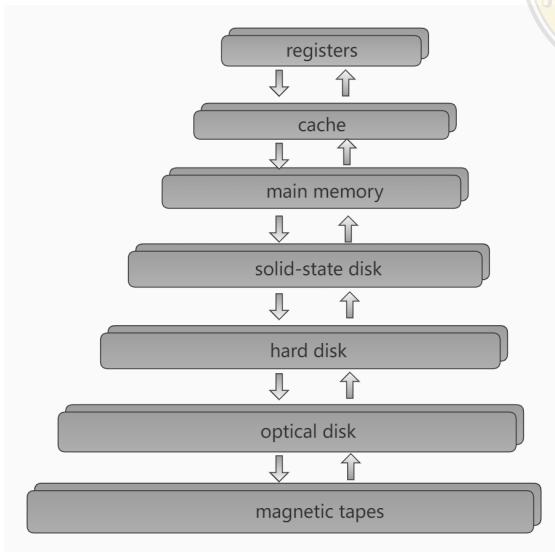
Interrupt Handling



- The operating system preserves the state of the CPU by storing registers and the program counter.
- Determines which type of interrupt has occurred:
 - polling, or
 - vectored interrupt system
- Separate segments of code determine what action should be taken for each type of interrupt.

Computer-System Organization

Storage-Device Hierarchy





Computer-System Organization

Storage Structure

- Main memory — only large storage media that the CPU can access directly.
 - Random access.
 - Typically volatile.
- Secondary storage — extension of main memory that provides large nonvolatile storage capacity.
- Hard disks — rigid metal or glass platters covered with magnetic recording material.
 - Disk surface is logically divided into tracks, which are subdivided into sectors.
 - The disk controller determines the logical interaction between the device and the computer.
- Solid-state disks — faster than hard disks, nonvolatile.
 - Various technologies.
 - Becoming more popular.



Computer-System Organization

I/O Structure — No Interrupts

- After I/O starts, control returns to user program only upon I/O completion
 - Wait instruction idles the CPU.
 - Wait loop (contention for memory access).
 - At most one I/O request is outstanding at a time, no simultaneous I/O processing.



Computer-System Organization

I/O Structure — Interrupt Driven

- After I/O starts, control returns to user program without waiting for I/O completion
 - **System call** — request to the OS to allow user to wait for I/O completion.
 - **Device-status table** contains entry for each I/O device indicating its type, address, and state.
 - OS indexes into I/O device table to determine device status and to modify table entry to include interrupt.
- Who returns the control?
 - **Device Driver** for each device controller to manage I/O. Provides uniform interface between controller and kernel.



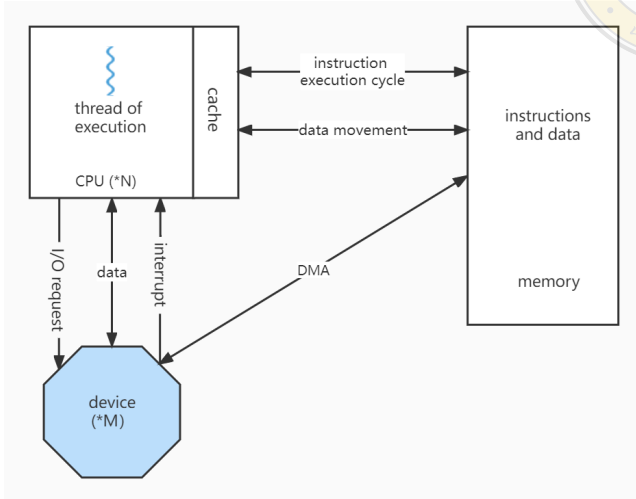
Computer-System Organization

I/O Structure — Direct Memory Access

- **Direct Memory Access** Structure
 - Used for high-speed I/O devices able to transmit information at close to memory speeds.
 - Device controller transfers blocks of data from buffer storage directly to main memory without CPU intervention.
 - Only one interrupt is generated per block, rather than the one interrupt per byte.

Computer-System Organization

How a Modern Computer Works





Contents

1. What Operating Systems Do
2. Computer-System Organization
- 3. Computer-System Architecture**
4. Operating-System Structure
5. Operating-System Operations
6. Operating-System Functions
7. Kernel Data Structures
8. Open-Source Operating Systems



Computer-System Architecture

- Most systems use a single general-purpose processor until 10 years ago.
 - Most systems have special-purpose processors as well.
- **Multiprocessor** systems growing in use and importance.
 - Also known as **parallel systems**, **tightly-coupled systems**.
 - Advantages include:
 1. Increased throughput.
 2. Economy of scale.
 3. Increased reliability — graceful degradation or fault tolerance.
 - Two types:
 1. **Asymmetric Multiprocessing** — each processor is assigned a specific task.
 2. **Symmetric Multiprocessing** — each processor performs all tasks.



Computer-System Architecture

A Dual-Core Design

- Multi-chip and **multicore**.
 - Why multicore?
 - On-chip communication is faster, uses significantly less power.
- Blade Servers: all in one chassis.
 - Chassis containing multiple separate systems.



Computer-System Architecture

Clustered Systems

- Like multiprocessor systems, but multiple systems working together.
- Loosely-coupled systems.
 - Usually sharing storage via a storage-area network (SAN).
 - Provides a **high-availability** service which survives failures.
 - » Asymmetric clustering has one machine in hot-standby mode.
 - » Symmetric clustering has multiple nodes running applications, monitoring each other.
 - Some clusters are for **high-performance** computing (HPC).
 - » Applications must be written to use parallelization.
 - Some have distributed lock manager (DLM) to avoid conflicting operations.



Contents

1. What Operating Systems Do
2. Computer-System Organization
3. Computer-System Architecture
- 4. Operating-System Structure**
5. Operating-System Operations
6. Operating-System Functions
7. Kernel Data Structures
8. Open-Source Operating Systems



Operating-System Structure


Some History

- No operating system. (1946 ~ 1955, ENIAC)
- A typical simple **Batch system** (since 1950s)
 - bring cards to IBM 1401
 - IBM 1401 read cards to tape
 - put tape on IBM 7094 which does computing
 - put tape on IBM 1401 which prints output



Operating-System Structure

Some History (contd.)



- **Multiprogramming** (Batch system) needed for efficiency. (since 1960s)
 - Single user cannot keep CPU and I/O devices busy at all times.
 - Multiprogramming organizes jobs (code and data) so CPU always has one to execute.
 - A subset of total jobs in system is kept in memory.
 - One job selected and run via **job scheduling (long term scheduling)**.
 - When it has to wait (for I/O for example), OS switches to another job.



Operating-System Structure

Some History (contd.)

- **Timesharing** (multitasking, since 1960s) is logical extension in which CPU switches jobs so frequently that users can interact with each job while it is running, creating **interactive** computing.
 - **Response time** should be < 1 second.
 - Each user has at least one program executing in memory — **process**.
 - If several jobs ready to run at the same time — **CPU scheduling (short term scheduling)**.
 - If processes don't fit in memory, **swapping** moves them in and out to run.
 - **Virtual memory** allows execution of processes not completely in memory.



In Class Exercise

Multiprogramming

- Consider a multiprogramming environment, two programs A and B share the system simultaneously, and run as follows.
 - For every 35 min, A runs on CPU for 15 min, then waits for I/O for 20 min.
 - For every 20 min, B runs on CPU for 10 min, then waits for I/O for 10 min.

Suppose B runs first, I/O_A and I/O_B are different devices, and the time of switching between A and B can be ignored. Draw the timeline for these two programs, and calculate CPU utilization within 60 minutes.

- What if I/O_A and I/O_B are the same device?




In Class Exercise


Key to Multiprogramming

Key to Q1: CPU utilization = $50/60 = 83.33\%$

v1:



v2:






In Class Exercise

Key to Multiprogramming

Key to Q2: $CPU\ utilization = 50/60 = 83.33\%$





Contents

1. What Operating Systems Do
2. Computer-System Organization
3. Computer-System Architecture
4. Operating-System Structure
- 5. Operating-System Operations**
6. Operating-System Functions
7. Kernel Data Structures
8. Open-Source Operating Systems



Operating-System Operations

Interrupt Driven (hardware and software)

- Hardware interrupt by one of the devices.
- Software interrupt (**exception** or **trap**):
 - Software error (e.g., division by zero).
 - Request for operating system service.
 - Other process problems include infinite loop, processes modifying each other or the operating system.



Operating-System Operations

Dual Mode

- **Dual-mode** operation allows OS to protect itself and other system components.
 - **User mode** and **kernel mode**.
 - **Mode bit** provided by hardware.
 - » Provides ability to distinguish when system is running user code or kernel code.
 - » Some instructions designated as **privileged**, only executable in kernel mode.
 - » **System call** changes mode to kernel, return from call resets it to user.
- Increasingly CPUs support multi-mode operations.
 - i.e., **virtual machine manager (VMM)** mode for guest **VMs**.

Operating-System Operations

Timer



- **Timer** to prevent infinite loop / process hogging resources.
 - Timer is set to interrupt the computer after some time period.
 - Keep a counter that is decremented by the physical clock.
 - Operating system set the counter (privileged instruction).
 - When counter zero generate an interrupt.
 - Set up before scheduling process to regain control or terminate program that exceeds allotted time.



Operating-System Operations

What Happens When the Computer System is Booting?

OS @boot (kernel mode)

initialize trap table

Hardware

remember addresses of ...
system call handler
timer handler
illegal instruction handler

start interrupt timer

start timer; interrupt after X ms



Operating-System Operations

What Happens When the Computer System is Running?

OS @run (kernel mode)	Hardware	Program (user mode)
to start process A: <i>return-from-trap</i> (into A)		
	move to <i>user mode</i>	
		process A runs: fetch instruction execute instruction ...
	timer interrupt move to <i>kernel mode</i> jump to interrupt handler	



Contents

1. What Operating Systems Do
2. Computer-System Organization
3. Computer-System Architecture
4. Operating-System Structure
5. Operating-System Operations
- 6. Operating-System Functions**
7. Kernel Data Structures
8. Open-Source Operating Systems



Operating-System Functions

- User interface, in Chapter 2
- CPU management (or **Process management**), in Chapter 3 ~ 7
 - Interrupt handling and management
 - CPU scheduling
 - Context switch (save & restore)
- **Memory management**, in Chapter 8 ~ 9
 - Physical memory allocation and deallocation
 - Address translation (logical \leftrightarrow physical)
 - Memory virtualization
- **Storage management**, in Chapter 10 ~ 12
 - Storage space allocation
 - File & directory management
- Device management, in Chapter 13
- Protection and security, in Chapter 14 ~ 15

Process/Memory/Storage Management

Too Many Terms?



- Just remember three key ideas for now:
 - **Virtualization**: The OS takes a physical resource (such as the processor, or memory, or a disk) and transforms it into a more general, powerful, and easy-to-use virtual form of itself.
 - **Concurrency**: Many problems arise, and must be addressed, when working on many things at once (i.e., concurrently) in the same program.
 - **Persistence**: The OS makes information persist, despite computer crashes, disk failures, or power outages.



Contents

1. What Operating Systems Do
2. Computer-System Organization
3. Computer-System Architecture
4. Operating-System Structure
5. Operating-System Operations
6. Operating-System Functions
- 7. Kernel Data Structures**
8. Open-Source Operating Systems

Kernel Data Structures

You Should Have Known the Concepts of ...



- arrays, lists, stacks, and queues,
- trees,
- hash functions and maps,
- bitmaps.



Contents

1. What Operating Systems Do
2. Computer-System Organization
3. Computer-System Architecture
4. Operating-System Structure
5. Operating-System Operations
6. Operating-System Functions
7. Kernel Data Structures
8. Open-Source Operating Systems



Open-Source Operating Systems

- Operating systems made available in source-code format rather than just binary closed-source.
- Counter to the copy protection and Digital Rights Management (DRM) movement.
- Started by Free Software Foundation (FSF), which has “copyleft” GNU Public License (GPL).
- Examples include GNU/Linux and BSD UNIX (including core of Mac OS X), and many more.
- Can use VMM like VMware Player (Free on Windows), Virtualbox (open source and free on many platforms - available @ <http://www.virtualbox.com>) .
 - Use to run guest operating systems for exploration.

openEuler Operating System



- Xv6
 - Xv6 documentation (PDF): [link](#)
 - Xv6 codes (PDF): [link](#)
 - Xv6 structure (mind): [link](#)
- EulerOS
 - A commercial Linux distribution developed by Huawei based on CentOS source code for enterprise applications.
 - Available @ <https://openeuler.org/en/>